

# BCG- TASK 2

July 07, 2023

## 1 Libraries

```
[350]: # Data analysis and wrangling
import pandas as pd
import numpy as np

# Data visualisation
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Dates
import datetime

#remove warnings
import warnings
warnings.filterwarnings('ignore')
```

## 2 Import and read data

```
[351]: churn_data=pd.read_csv('ml_case_training_output.csv')
churn_data.head()
```

```
[351]:
```

	id	churn
0	48ada52261e7cf58715202705a0451c9	0
1	24011ae4ebbe3035111d65fa7c15bc57	1
2	d29c2c54acc38ff3c0614d0a653813dd	0
3	764c75f661154dac3a6c254cd082ea7d	0
4	bba03439a292a1e166f80264c16191cb	0

```
[352]: churn_data['churn'] =
churn_data['churn'].replace({0:'stayed',1:'churned'})
churn_data.head()
```

```
[352]:
```

	id	churn
0	48ada52261e7cf58715202705a0451c9	stayed
1	24011ae4ebbe3035111d65fa7c15bc57	churned
2	d29c2c54acc38ff3c0614d0a653813dd	stayed
3	764c75f661154dac3a6c254cd082ea7d	stayed
4	bba03439a292a1e166f80264c16191cb	stayed

```
[353]: history_data=pd.read_csv('ml_case_training_hist_data.csv')
history_data.head()
```

```
[353]:
```

	id	price_date	price_p1_var	price_p2_var \
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	0.0
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	0.0
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	0.0
3	038af19179925da21a25619c5a24b745	2015-04-01	0.149626	0.0
4	038af19179925da21a25619c5a24b745	2015-05-01	0.149626	0.0

	price_p3_var	price_p1_fix	price_p2_fix	price_p3_fix
0	0.0	44.266931	0.0	0.0
	0.0	0.0	20.0	44.266931
	0.0	0.0	0.0	30.0
	44.266931	0.0	0.0	
4	0.0	44.266931	0.0	0.0

```
[354]: training_data=pd.read_csv('ml_case_training_data.csv')
training_data.tail()
```

```
[354]:
```

	id	activity_new	campaign_disc_ele \
16091	18463073fb097fc0ac5d3e040f356987	NaN	NaN
16092	d0a6f71671571ed83b2645d23af6de00	NaN	NaN
16093	10e6828ddd62cbcf687cb74928c4c2d2	NaN	NaN
16094	1cf20fd6206d7678d5bcafd28c53b4db	NaN	NaN
16095	563dde550fd624d7352f3de77c0cdfcd	NaN	NaN

	channel_sales	cons_12m	cons_gas_12m \
16091	foosdfpfkusacimwkcsosbicdxkica	32270	47940
16092	foosdfpfkusacimwkcsosbicdxkica	7223	0
16093	foosdfpfkusacimwkcsosbicdxkica	1844	0
16094	foosdfpfkusacimwkcsosbicdxkica	131	0
16095	NaN	8730	0

	cons_last_month	date_activ	date_end	date_first_activ ... \
16091	0	2012-05-24	2016-05-08	NaN ...
16092	181	2012-08-27	2016-08-27	2012-08-27 ...
16093	179	2012-02-08	2016-02-07	NaN ...
16094	0	2012-08-30	2016-08-30	NaN ...
16095	0	2009-12-18	2016-12-17	NaN ...

	forecast_price_pow_p1	has_gas	imp_cons	margin_gross_pow_ele \
16091	44.311378	t	0.00	27.88
16092	58.995952	f	15.94	0.00
16093	40.606701	f	18.05	39.84

```

16094      44.311378      f      0.00      13.08
16095      45.311378      f      0.00      11.84
      margin_net_pow_ele nb_prod_act net_margin num_years_antig \
16091      27.88      2      381.77      4
16092      0.00      1      90.34 3
16093      39.84      1      20.38 4
16094      13.08      1      0.96 3
16095      11.84      1      96.34 6

```

origin\_up pow\_max

```

16091 lxidpiddsbxsbosboudacockeimpuepw 15.000
16092 lxidpiddsbxsbosboudacockeimpuepw6.000
16093 lxidpiddsbxsbosboudacockeimpuepw15.935
16094 lxidpiddsbxsbosboudacockeimpuepw11.000
16095 ldkssxwpmemidmecebumciepifcamkci10.392

```

[5 rows x 32 columns]

```

[355]: merge=pd.merge(churn_data, training_data, on='id')
merge.tail()

```

```

[355]:      id churn activity_new \
16091 18463073fb097fc0ac5d3e040f356987stayed      NaN
16092 d0a6f71671571ed83b2645d23af6de00 churned      NaN
16093 10e6828ddd62cbcf687cb74928c4c2d2 churned      NaN
16094 1cf20fd6206d7678d5bcafd28c53b4dbstayed      NaN
16095 563dde550fd624d7352f3de77c0cdfcdstayed      NaN
      campaign_disc_ele      channel_sales cons_12m \
16091      NaN foosdfpfkusacimwkcsosbicdxkicaua 32270
16092      NaN foosdfpfkusacimwkcsosbicdxkicaua 7223
16093      NaN foosdfpfkusacimwkcsosbicdxkicaua 1844
16094      NaN foosdfpfkusacimwkcsosbicdxkicaua 131
16095      NaN      NaN      NaN 8730
      cons_gas_12m cons_last_month date_activ date_end ... \
16091      47940      0 2012-05-24 2016-05-08
      ...
16092      0      181 2012-08-27 2016-08-27
      ...
16093      0      179 2012-02-08 2016-02-07
      ...
16094      0      0 2012-08-30 2016-08-30
      ...
16095      0      0 2009-12-18 2016-12-17 ...
forecast_price_pow_p1 has_gas imp_cons margin_gross_pow_ele \

```

```

16091      44.311378  t      0.00 27.88
16092      58.995952  f      15.94 0.00
16093      40.606701      f      18.05      39.84
16094      44.311378      f      0.00      13.08
16095      45.311378      f      0.00      11.84
      margin_net_pow_ele nb_prod_act net_margin num_years_antig \
16091      27.88      2      381.77      4
16092      0.00      1      90.34 3
16093      39.84      1      20.38 4
16094      13.08      1      0.96 3
16095      11.84      1      96.34 6
origin_up pow_max
16091 lxicpiddsbxsbosboudacockeimpuepw 15.000
16092 lxicpiddsbxsbosboudacockeimpuepw6.000
16093 lxicpiddsbxsbosboudacockeimpuepw15.935
16094 lxicpiddsbxsbosboudacockeimpuepw11.000
16095 ldkssxwpmemidmecebumciepifcamkci10.392
[5 rows x 33 columns]

```

### 3 CHURN DATA

```

[356]: churn_data.count()

[356]: id    16096 churn
      16096
      dtype: int64

[357]: churn_count=churn_data['churn'].value_counts()
      print(churn_count)

      stayed    14501
      churned    1595
      Name: churn, dtype: int64

```

It can be seen that, the number of companies that have churned out is

```

[358]: rate_of_churn = pd.DataFrame(churn_data['churn'].value_counts() /
churn_data.
      .shape[0] * 100)
      print(rate_of_churn )

      churn

```

```

stayed      90.090706
churned     9.909294

```

It can be seen that, the number of companies that have churned out is 1595 which represent 9.9%, approximately, 10%.

```

[359]: #changing the column names
merge['churn'] = merge['churn'].replace({0:'stayed',1:'churned'})
merge.head()

```

```

[359]:
      id      churn \
0  48ada52261e7cf58715202705a0451c9    stayed
1  24011ae4ebbe3035111d65fa7c15bc57  churned
2  d29c2c54acc38ff3c0614d0a653813dd    stayed
3  764c75f661154dac3a6c254cd082ea7d    stayed
4  bba03439a292a1e166f80264c16191cb    stayed

      activity_new campaign_disc_ele \
0  esoiifxdlbkcsluxmfuacbdckommixw      NaN
1                                NaN      NaN
2                                NaN      NaN
3                                NaN      NaN
4                                NaN      NaN

      channel_sales cons_12m cons_gas_12m cons_last_month \
0  lmkebamcaaclubfxadlmueccxoimlema309275      0      10025
1  foosdfpfkusacimwkcsoibcdxkicaua      0      54946      0
2                                NaN      4660      0      0
3  foosdfpfkusacimwkcsoibcdxkicaua      544      0      0
4  lmkebamcaaclubfxadlmueccxoimlema      1584      0      0

      date_activdate_end ... forecast_price_pow_p1 has_gas imp_cons \
0  2012-11-07 2016-11-06 ... 58.995952 f 831.8 1 2013-06-15
2016-06-15 ... 40.606701 t 0.0
2  2009-08-21 2016-08-30 ... 44.311378 f      0.0
3  2010-04-16 2016-04-16 ... 44.311378 f      0.0
4  2010-03-30 2016-03-30 ... 44.311378 f      0.0

      margin_gross_pow_ele margin_net_pow_ele nb_prod_act net_margin \
0      -41.76-41.76      1      1732.36
1      25.44 25.44 2      678.99
2      16.38      16.38      1      18.89
3      28.60      28.60      1      6.60

```

4	30.22	30.22	1	25.46
---	-------	-------	---	-------

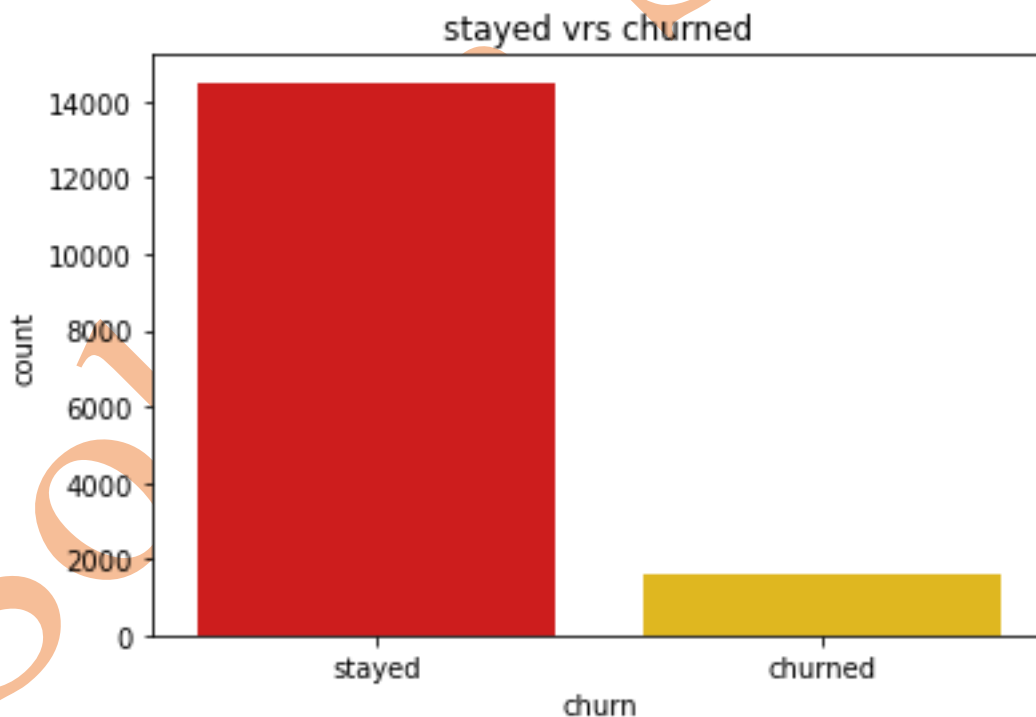
	num_years_antig	origin_up	pow_max
0	3 ldkssxwpmemidmecebumciepifcamkci		180.000
1	3 lxidpiddsbxsbosboudacockeimpuepw		43.648
2	6 kamkkxfxxuwbdslkwifmmcsiusuosws		13.800
3	6 kamkkxfxxuwbdslkwifmmcsiusuosws		13.856
4	6 kamkkxfxxuwbdslkwifmmcsiusuosws		13.200

[5 rows x 33 columns]

## 4 Data visualization of churn

```
[360]: sns.countplot(x= 'churn', data = churn_data, palette = 'hot')
plt.title('stayed vrs churned')
```

```
[360]: Text(0.5, 1.0, 'stayed vrs churned')
```



## 5 Describing data

```
[361]: merge.info()
```

```

<class
'pandas.core.frame.DataFrame'>
Int64Index: 16096 entries, 0 to
16095 Data columns (total 33
columns):
#      Column                                Non-Null Count
-----
0      id                                16096 non-null
                                object
1      churn                            16096 non-null
                                object
2      activity_new                      6551 non-null object
3      campaign_disc_ele                 0 non-null    float64
4      channel_sales                    11878 non-null
                                object
5      cons_12m                         16096 non-null
                                int64
6      cons_gas_12m                    16096 non-null
                                int64
7      cons_last_month                 16096 non-null
                                int64
8      date_activ                      16096 non-null
                                object
9      date_end                        16094 non-null
                                object
10     date_first_activ                3508 non-null object
11     date_modif_prod                15939 non-null
                                object
12     date_renewal                   16056 non-null
                                object
13     forecast_base_bill_ele         3508 non-null
                                float64
14     forecast_base_bill_year        3508 non-null
                                float64
15     forecast_bill_12m              3508 non-null
                                float64
16     forecast_cons                  3508 non-null
                                float64
17     forecast_cons_12m              16096 non-null
                                float64
18     forecast_cons_year             16096 non-null
                                int64
19     forecast_discount_energy       15970 non-null float64
20     forecast_meter_rent_12m        16096 non-null float64
21     forecast_price_energy_p1       15970 non-null float64

```

```

22 forecast_price_energy_p2 15970 non-null float64
23 forecast_price_pow_p1    15970      non-null
    float64
24 has_gas                  16096 non-null
    object
25 imp_cons                 16096      non-null
    float64
26 margin_gross_pow_ele    16083      non-null
    float64
27 margin_net_pow_ele      16083      non-null
    float64
28 nb_prod_act             16096 non-null
    int64
29 net_margin              16081      non-null
    float64
30 num_years_antig        16096 non-null
    int64
31 origin_up              16009 non-null
    object
32 pow_max                16093      non-null
    float64

```

```

dtypes: float64(16), int64(6),
object(11) memory usage: 4.2+ MB

```

It can be seen that the types of date is object, but needs to be in datetime.

```
[362]: merge.describe()
```

```

[362]:   campaign_disc_ele  cons_12m cons_gas_12m cons_last_month \
count          0.0  1.609600e+04          1.609600e+04
      1.609600e+04
mean          NaN  1.948044e+05          1.946154e+04
      3.191164e+04
std          NaN  6.795151e+05          8.235676e+04
      1.775885e+05
min          NaN  -1.252760e+05  -
      3.037000e+03          9.138600e+04
25%          NaN  5.906250e+03          0.000000e+00
      0.000000e+00
50%          NaN  1.533250e+04          9.010000e+02
      0.000000e+00
75%          NaN  5.022150e+04          4.127000e+03
      0.000000e+00
max          NaN  1.609711e+07          4.538720e+06
      4.188440e+06
forecast_base_bill_ele forecast_base_bill_year forecast_bill_12m \
count          3508.000000          3508.000000          3508.000000
mean           335.843857          335.843857          3837.441866
std            649.406000          649.406000          5425.744327

```



min	-364.940000	-364.940000	-
			2503.480000
25%	0.000000	0.000000	1158.175000
50%	162.955000	162.955000	2187.230000
75%	396.185000	396.185000	4246.555000
max	12566.080000	12566.080000	81122.630000

	forecast_cons	forecast_cons_12m	forecast_cons_year ... \
count	3508.000000	16096.000000	16096.000000 ...
mean	206.845165	2370.555949	1907.347229 ...
std	455.634288	4035.085664	5257.364759 ...
min	0.000000	-16689.260000	-85627.000000 ...
25%	0.000000	513.230000	0.000000 ...
50%	42.215000	1179.160000	378.000000 ...
75%	228.117500	2692.077500	1994.250000 ...
max	9682.890000	103801.930000	175375.000000 ...

	forecast_price_energy_p1	forecast_price_energy_p2 \
count	15970.000000	15970.000000
mean	0.135901	0.052951
std	0.026252	0.048617
min	0.000000	0.000000
25%	0.115237	0.000000
50%	0.142881	0.086163
75%	0.146348	0.098837
max	0.273963	0.195975

	forecast_price_pow_p1	imp_cons	margin_gross_pow_ele \
count	15970.000000	16096.000000	16083.000000
mean	43.533496	196.123447	22.462276
std	5.212252	494.366979	23.700883
min	-0.122184	-9038.210000	-525.540000
25%	40.606701	0.000000	11.960000
50%	44.311378	44.465000	21.090000
75%	44.311378	218.090000	29.640000
max	59.444710	15042.790000	374.640000

	margin_net_pow_elenb	prod_actnet	margin_net	num_years_antig \
count	16083.000000	16096.000000	16081.000000	16096.000000
mean	21.460318	1.347788	217.987028	5.030629
std	27.917349	1.459808	366.742030	1.676101
min	-615.660000	1.000000	-4148.990000	1.000000
25%	11.950000	1.000000	51.970000	4.000000
50%	20.970000	1.000000	119.680000	5.000000
75%	29.640000	1.000000	275.810000	6.000000

```
max          374.640000    32.000000 24570.650000    16.000000
```

```
      pow_max
count 16093.000000
mean    20.604131
std     21.772421
min      1.000000
25%     12.500000
50%     13.856000
75%     19.800000
max     500.000000
[8 rows x 22 columns]
```

```
[363]: merge['has_gas'] = merge['has_gas'].replace({'f':'No','t':'Yes'})
merge.head()
```

```
[363]:
      id      churn \
0  48ada52261e7cf58715202705a0451c9  stayed
1  24011ae4ebbe3035111d65fa7c15bc57  churned
2  d29c2c54acc38ff3c0614d0a653813dd  stayed
3  764c75f661154dac3a6c254cd082ea7d  stayed
4  bba03439a292a1e166f80264c16191cb  stayed

      activity_new campaign_disc_ele \
0  esoiifxdlbkcsluxmfuacbdckommixw  NaN
1                                     NaN  NaN
2                                     NaN  NaN
3                                     NaN  NaN
4                                     NaN  NaN

      channel_sales cons_12m cons_gas_12m cons_last_month \
0  lmkebamcaaclubfxadlmueccxoimlema309275      0      10025
1  foosdfpfkusacimwkcsosbicdxkicaua      0      54946      0
2                                     NaN      4660      0      0
3  foosdfpfkusacimwkcsosbicdxkicaua      544      0      0
4  lmkebamcaaclubfxadlmueccxoimlema      1584      0      0

      date_activdate_end ... forecast_price_pow_p1 has_gas imp_cons \
0  2012-11-07 2016-11-06 ... 58.995952 No      831.8
1  2013-06-15 2016-06-15 ... 40.606701 Yes      0.0
2  2009-08-21 2016-08-30 ... 44.311378 No      0.0
3  2010-04-16 2016-04-16 ... 44.311378 No      0.0
4  2010-03-30 2016-03-30 ... 44.311378 No      0.0

      margin_gross_pow_ele margin_net_pow_ele nb_prod_act net_margin \
0      -41.76      -41.76      1      1732.36
1      25.44      25.44      2      678.99
2      16.38      16.38      1      18.89
3      28.60      28.60      1      6.60
```

```

4          30.22          30.22          1          25.46
   num_years_antig          origin_up pow_max
0          3 ldkssxwpmemidmecebumciepifcamkci 180.000
1          3 lxidpiddsbxsbsboudacockeimpuepw          43.648
2          6 kamkkxfxxuwbdslkwifmmcsiusiusws          13.800
3          6 kamkkxfxxuwbdslkwifmmcsiusiusws          13.856
4          6 kamkkxfxxuwbdslkwifmmcsiusiusws          13.200

```

[5 rows x 33 columns]

[364]: history\_data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               193002 non-null object
1   price_date       193002 non-null object
2   price_p1_var     191643 non-null float64
3   price_p2_var     191643 non-null float64
4   price_p3_var     191643 non-null float64
5   price_p1_fix     191643 non-null float64
6   price_p2_fix     191643 non-null float64
7   price_p3_fix     191643 non-null float64
dtypes: float64(6), object(2)
memory usage: 11.8+ MB

```

[365]: history\_data.describe()

```

[365]: price_p1_var price_p2_var price_p3_var price_p1_fix \ count
191643.000000 191643.000000 191643.000000 191643.000000
mean          0.140991      0.054412      0.030712      43.325546
std           0.025117      0.050033      0.036335       5.437952
min           0.000000      0.000000      0.000000     -0.177779
25%           0.125976      0.000000      0.000000     40.728885
50%           0.146033      0.085483      0.000000     44.266930
75%           0.151635      0.101780      0.072558     44.444710
max           0.280700      0.229788      0.114102     59.444710

           price_p2_fix price_p3_fix
count 191643.000000
191643.000000 mean      10.698201
              std        6.455436
              min        7.782279
              25%       -0.065172
              0.000000

```

```
50%    0.000000    0.000000    75%
24.339581 16.226389
```

```
max      36.490692    17.458221
```

## 6 Checking for missing data

```
[366]: missing_figures_1 = history_data.isnull().sum()
missing_figures_1 = missing_figures_1[missing_figures_1
> 0]
pd.DataFrame({"missing_figures_1": missing_figures_1,
"Missing_values_1(%)":
→history_data.isnull().sum()/len(history_data.index)*100}).sort_va
lues(by =
→"Missing_values_1(%)", ascending = False)
```

```
[366]: missing_figures_1 Missing_values_1(%)
price_p1_fix      1359.0      0.704138
price_p1_var      1359.0      0.704138
price_p2_fix      1359.0      0.704138
price_p2_var      1359.0      0.704138
price_p3_fix      1359.0      0.704138
price_p3_var      1359.0      0.704138 id      NaN
0.000000 price_date      NaN      0.000000
```

```
[367]: missing_figures = merge.isnull().sum() missing_figures =
missing_figures[missing_figures > 0]
pd.DataFrame({"missing_figures": missing_figures, "Missing values
(%)": merge.
→isnull().sum()/len(merge.index)*100}).sort_values(by = "Missing
values (%)",
→ascending = False)
```

```
[367]: missing_figures Missing values (%) campaign_disc_ele
16096.0      100.000000 date_first_activ12588.0
78.205765 forecast_base_bill_ele 12588.0      78.205765
forecast_cons      12588.0      78.205765 forecast_bill_12m
12588.0      78.205765 forecast_base_bill_year
12588.0      78.205765 activity_new      9545.0
```

```

59.300447 channel_sales      4218.0      26.205268
date_modif_prod 157.0 0.975398 forecast_price_pow_p1  126.0
0.782803 forecast_price_energy_p2      126.0 0.782803
forecast_discount_energy  126.0 0.782803
forecast_price_energy_p1  126.0 0.782803 origin_up      87.0
0.540507 date_renewal 40.0 0.248509 net_margin  15.0
0.093191 margin_gross_pow_ele  13.0 0.080765
margin_net_pow_ele  13.0 0.080765 pow_max 3.0 0.018638
date_end      2.0      0.012425
forecast_meter_rent_12m      NaN      0.000000
forecast_cons_year      NaN      0.000000
date_activ      NaN      0.000000
has_gas      NaN      0.000000
id      NaN      0.000000
imp_cons      NaN      0.000000
cons_last_month      NaN      0.000000
cons_gas_12m      NaN      0.000000
nb_prod_act      NaN      0.000000
cons_12m      NaN      0.000000
num_years_antig      NaN      0.000000
churn      NaN      0.000000
forecast_cons_12m      NaN      0.000000

```

We need to drop columns with a lot of missing data. Hence all columns with more than 70% missing data should be dropped

The history data looks good with less than 1% missing data.

## 7 Visualization of missing figures

```


[368]: plt.figure(figsize=(20, 10))      cmap      =
sns.cubehelix_palette(light=1,      as_cmap=True,
reverse=True) sns.heatmap(merge.isnull(), cmap=cmap)

```

```

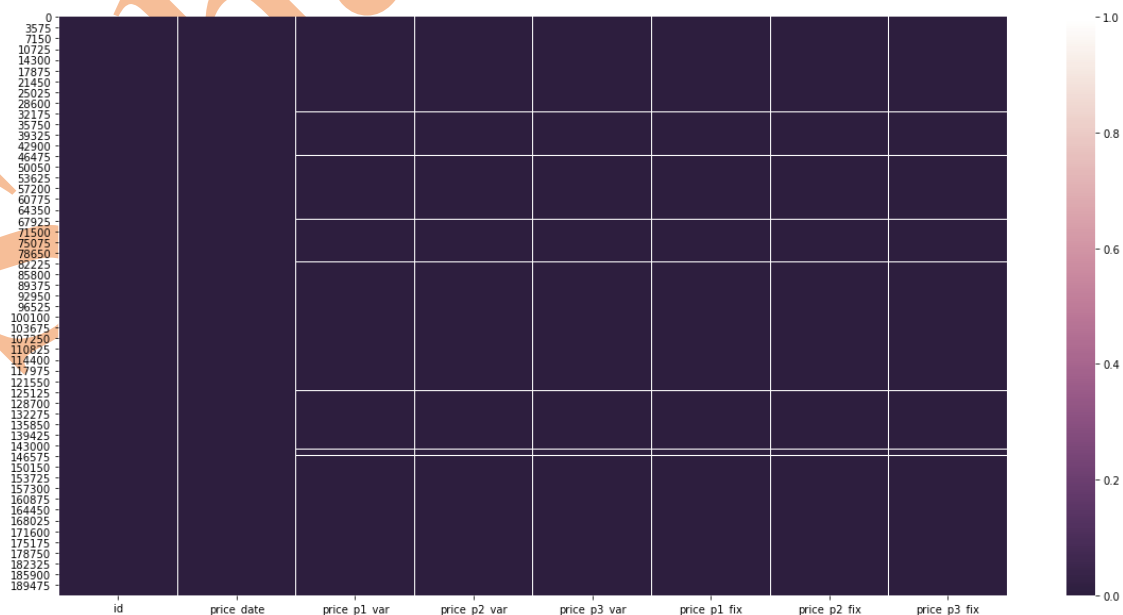
[368]: <AxesSubplot:>

```



```
[369]: plt.figure(figsize=(20, 10))
sns.cubehelix_palette(light=1, as_cmap=True, reverse=True)
sns.heatmap(history_data.isnull(),
cmap=cmap)
```

[369]: <AxesSubplot:>



```
[370]: pip install missingno
```

```
Requirement already satisfied: missingno in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (0.4.2)
Requirement already satisfied: numpy in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
missingno) (1.19.5)
Requirement already satisfied: matplotlib in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
missingno) (3.3.2)
Requirement already satisfied: scipy in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
missingno) (1.5.2)
Requirement already satisfied: seaborn in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
missingno) (0.11.0)
Requirement already satisfied:
pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
matplotlib->missingno) (2.4.7)
Requirement already satisfied: certifi>=2020.06.20 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
matplotlib->missingno) (2020.6.20)
Requirement already satisfied: python-dateutil>=2.1 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
matplotlib->missingno) (2.8.1)
Requirement already satisfied: pillow>=6.2.0 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
matplotlib->missingno) (8.0.1)
Requirement already satisfied: kiwisolver>=1.0.1 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
matplotlib->missingno) (1.3.0)
Requirement already satisfied: cycycler>=0.10 in
/Users/barbarazen/anaconda3/lib/python3.8/site-
packages (from matplotlib->missingno) (0.10.0)
Requirement already satisfied: six in
/Users/barbarazen/anaconda3/lib/python3.8/site-
packages (from cycycler>=0.10->matplotlib->missingno)
(1.15.0) Requirement already satisfied: pandas>=0.23
in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
seaborn->missingno) (1.1.3)
Requirement already satisfied: pytz>=2017.2 in
/Users/barbarazen/anaconda3/lib/python3.8/site-packages (from
pandas>=0.23->seaborn->missingno) (2020.1)
Note: you may need to restart the kernel to use updated packages.
```

```
[371]: import missingno as msno
```

The missingno correlation heatmap measures nullity correlation: how strongly the presence or absence of one variable affects the presence of another:

Nullity correlation ranges from -1 (if one variable appears the other definitely does not) to 0 (variables appearing or not appearing have no effect on one another) to 1 (if one variable appears the other definitely also does).

Variables that are always full or always empty have no meaningful correlation, and so are silently removed from the visualization—in this case for instance the datetime and injury number columns, which are completely filled, are not included.

```
[372]: msno.heatmap(history_data)
```


```
[372]: <AxesSubplot:>
```



```
[373]: msno.heatmap(merge)
```

```
[373]: <AxesSubplot:>
```





Variables that are always full or always empty have no meaningful correlation, and so are silently removed from the visualization—in this case for instance the datetime and injury number columns, which are completely filled, are not included. It can be seen that, the columns with a lot of missing data is already dropped from the correlation heatmap.

We need to drop columns with a lot of missing data. Hence all columns with more than 70% missing data should be dropped

## 8 Dropping missing figures

```
[374]: merge=merge.drop(columns= ["forecast_base_bill_ele",
"date_first_activ",
,"campaign_disc_ele", "forecast_base_bill_year", "forecast_bill_12m",
,"forecast_cons", ])
merge.head()
```

```
[374]:
```

	id	churn \
0	48ada52261e7cf58715202705a0451c9	stayed
1	24011ae4ebbe3035111d65fa7c15bc57	churned
2	d29c2c54acc38ff3c0614d0a653813dd	stayed
3	764c75f661154dac3a6c254cd082ea7d	stayed
4	bba03439a292a1e166f80264c16191cb	stayed

```

        activity_new          channel_sales \
0          esoiifxdlbkcsluxmfuacbdckommixw
          lmkebamcaclubfxadlmueccxoimlema
1          NaN foosdfpfkusacimwkcsosbicdxkicaua
2          NaN NaN
3          NaN foosdfpfkusacimwkcsosbicdxkicaua
4          NaN lmkebamcaclubfxadlmueccxoimlema

cons_12m cons_gas_12m cons_last_month date_activ date_end \
0          309275 0          10025 2012-11-07 2016-11-06
1          0 549460 2013-06-15 2016-06-15
2          4660 0          0 2009-08-21 2016-08-30
3          544 0 0 2010-04-16 2016-04-16 4 1584 0 0 2010-03-30 2016-
          03-30

date_modif_prod ... forecast_price_pow p1 has_gas imp_cons \
0          2012-11-07 ...          58.995952          No          831.8
1          NaN ...          40.606701          Yes          0.0
2          2009-08-21 ...          44.311378          No          0.0
3          2010-04-16 ...          44.311378          No          0.0
4          2010-03-30 ...          44.311378          No          0.0
margin_gross_pow_ele margin_net_pow_ele nb_prod_act net_margin \
0          -41.76-41.76          1          1732.36
1          25.44 25.44 2          678.99
2          16.38          16.38          1          18.89
3          28.60          28.60          1          6.60
4          30.22          30.22          1          25.46

num_years_antig          origin_up pow_max
0          3 ldkssxwpmemidmecebumciepifcamkci
          180.000
1          3 lxdpiddsbxsbosboudacockeimpuepw
          43.648
2          6 kamkkxfxxuwbdslkwifmmcsiusiusws
          13.800
3          6 kamkkxfxxuwbdslkwifmmcsiusiusws
          13.856
4          6 kamkkxfxxuwbdslkwifmmcsiusiusws
          13.200
[5 rows x 27 columns]

```

## 8.1 Replacing the Missing data with the mean of the data

To start, we need to find the mean and replace the mean with the null values

```

[375]: #finding the mean of the data mean_cons_12m=
merge["cons_12m"].mean() mean_cons_gas_12m=

```

```

merge["cons_gas_12m"].mean() mean_cons_last_month=
merge["cons_last_month"].mean() mean_forecast_cons_12m=
merge["forecast_cons_12m"].mean() mean_forecast_cons_year=
merge["forecast_cons_year"].mean()
mean_forecast_discount_energy=
merge["forecast_discount_energy"].mean()
mean_forecast_meter_rent_12m=
merge["forecast_meter_rent_12m"].mean()
mean_forecast_price_energy_p1=
merge["forecast_price_energy_p1"].mean()
mean_forecast_price_energy_p2=
merge["forecast_price_energy_p2"].mean()
mean_forecast_price_pow_p1=
merge["forecast_price_pow_p1"].mean() mean_imp_cons=
merge["imp_cons"].mean() mean_margin_gross_pow_ele=
merge["margin_gross_pow_ele"].mean() mean_margin_net_pow_ele=
merge["margin_net_pow_ele"].mean() mean_nb_prod_act=
merge["nb_prod_act"].mean() mean_net_margin=
merge["net_margin"].mean() mean_num_years_antig=
merge["num_years_antig"].mean() mean_pow_max=
merge["pow_max"].mean()

```

```

[376]: merge["cons_12m"] = merge["cons_12m"].fillna(mean_cons_12m)
merge["cons_gas_12m"] =
merge["cons_gas_12m"].fillna(mean_cons_gas_12m)
merge["cons_last_month"] =
merge["cons_last_month"].fillna(mean_cons_last_month)
merge["forecast_cons_12m"] = merge["forecast_cons_12m"].
    .fillna(mean_forecast_cons_12m)
merge["forecast_cons_year"] =
merge["forecast_cons_year"].
    .fillna(mean_forecast_cons_year)
merge["forecast_discount_energy"] =
merge["forecast_discount_energy"].
    .fillna(mean_forecast_discount_energy)
merge["forecast_meter_rent_12m"] =
merge["forecast_meter_rent_12m"].
    .fillna(mean_forecast_meter_rent_12m)
merge["forecast_price_energy_p1"] =
merge["forecast_price_energy_p1"].

```

```

    .fillna(mean_forecast_price_energy_p1)
merge["forecast_price_energy_p2"] =
merge["forecast_price_energy_p2"].
    .fillna(mean_forecast_price_energy_p2)
merge["forecast_price_pow_p1"] =
merge["forecast_price_pow_p1"].
    .fillna(mean_forecast_price_pow_p1) merge["imp_cons"]
= merge["imp_cons"].fillna(mean_imp_cons)
merge["margin_gross_pow_ele"] =
merge["margin_gross_pow_ele"].
    .fillna(mean_margin_gross_pow_ele)
merge["margin_net_pow_ele"] =
merge["margin_net_pow_ele"].
    .fillna(mean_margin_net_pow_ele) merge["nb_prod_act"] =
merge["nb_prod_act"].fillna(mean_nb_prod_act) merge["net_margin"] =
merge["net_margin"].fillna(mean_net_margin) merge["num_years_antig"]
= merge["num_years_antig"].fillna(mean_num_years_antig)
merge["pow_max"] = merge["pow_max"].fillna(mean_pow_max)

```

```

[377]: mean_price_p1_var=
history_data["price_p1_var"].mean()
mean_price_p2_var=
history_data["price_p2_var"].mean()
mean_price_p3_var=
history_data["price_p3_var"].mean()
mean_price_p1_fix=
history_data["price_p1_fix"].mean()
mean_price_p2_fix=
history_data["price_p2_fix"].mean()
mean_price_p3_fix=
history_data["price_p3_fix"].mean()

```

```

[378]: history_data["price_p1_var"] = history_data["price_p1_var"].
    .fillna(mean_price_p1_var)
history_data["price_p2_var"] =
history_data["price_p2_var"].
    .fillna(mean_price_p2_var)
history_data["price_p3_var"] =
history_data["price_p3_var"].

```

```

    .fillna(mean_price_p3_var)
history_data["price_p1_fix"] =
history_data["price_p1_fix"].
    .fillna(mean_price_p1_fix)
history_data["price_p2_fix"] =
history_data["price_p2_fix"].
    .fillna(mean_price_p2_fix)
history_data["price_p3_fix"] =
history_data["price_p3_fix"].
    .fillna(mean_price_p3_fix)

```

```
[379]: merge.isnull().sum()
```

```

[379]: id    0 churn      0 activity_new
      9545 channel_sales    4218
cons_12m    0 cons_gas_12m    0
cons_last_month 0 date_activ
      0 date_end 2
date_modif_prod 157 date_renewal
      40 forecast_cons_12m    0
forecast_cons_year    0
forecast_discount_energy    0
forecast_meter_rent_12m    0
forecast_price_energy_p1    0
forecast_price_energy_p2    0
forecast_price_pow_p10 has_gas
      0 imp_cons    0
margin_gross_pow_ele    0
margin_net_pow_ele    0
nb_prod_act    0 net_margin
      0 num_years_antig    0
origin_up 87 pow_max    0
dtype: int64

```


## 9 Checking Skewness

```

[380]: fig, axs = plt.subplots(figsize=(48,20))
      sns.boxplot(x=merge['cons_12m'])


```

```
[380]: <AxesSubplot:xlabel='cons_12m'>
```



```
[381]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['cons_last_month'])
```

```
[381]: <AxesSubplot:xlabel='cons_last_month'>
```




```
[382]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['cons_last_month'])
```

```
[382]: <AxesSubplot:xlabel='cons_last_month'>
```




```
[383]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['forecast_cons_12m'])
```

```
[383]: <AxesSubplot:xlabel='forecast_cons_12m'>
```




```
[384]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['forecast_cons_year'])
```

```
[384]: <AxesSubplot:xlabel='forecast_cons_year'>
```



```
[385]: fig, axs = plt.subplots(figsize=(28,10))
sns.boxplot(x=merge['imp_cons'])
```


```
[385]: <AxesSubplot:xlabel='imp_cons'>
```



```
[386]: fig, axs = plt.subplots(figsize=(28,10))
sns.boxplot(x=merge['forecast_discount_energy'])
```


```
[386]: <AxesSubplot:xlabel='forecast_discount_energy'>
```






```
[387]: fig, axs = plt.subplots(figsize=(28,10))
sns.boxplot(x=merge['forecast_meter_rent_12m'])
```

```
[387]: <AxesSubplot:xlabel='forecast_meter_rent_12m'>
```




```
[388]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['forecast_price_pow_p1'])
```

```
[388]: <AxesSubplot:xlabel='forecast_price_pow_p1'>
```




```
[389]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['margin_net_pow_ele'])
```

```
[389]: <AxesSubplot:xlabel='margin_net_pow_ele'>
```



```
[390]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['nb_prod_act'])
```

```
[390]: <AxesSubplot:xlabel='nb_prod_act'>
```




```
[391]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['net_margin'])
```

```
[391]: <AxesSubplot:xlabel='net_margin'>
```



```
[392]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['num_years_antig'])
```

```
[392]: <AxesSubplot:xlabel='num_years_antig'>
```



```
[393]: fig, axs = plt.subplots(figsize=(48,20))
sns.boxplot(x=merge['pow_max'])
```

```
[393]: <AxesSubplot:xlabel='pow_max'>
```



## 10 Finding factors that affect churning

Now let's find the code of the sales channel that companies subscribed to and check to see if it has correlation with the churning.


```
[394]: channelsales_count=pd.DataFrame(merge['channel_sales'].value_counts
()) print(channelsales_count)
```

channel\_sales

foosdfpfkusacimwkcsosbicdxkicaua	7377
lmkebamcaaclubfxadlmueccxoimlema	2073
usilxuppasemubllopkaafesmlibmsdf	1444
ewpakwlliwisiwduibdlfmalxowmwpci	966
sddiedcsflslkckwlfdpoeailfpeds	12
epumfxlbckeskwexbiuasklxalciuu	4
fixdbufsefwooaasfcxdxadsiekoeaa	2

```
[395]: pd.crosstab(merge['channel_sales'], merge['churn'], normalize='index')
        .plot.bar(stacked=True)
```

```
[395]: <AxesSubplot: xlabel='channel_sales'>
```

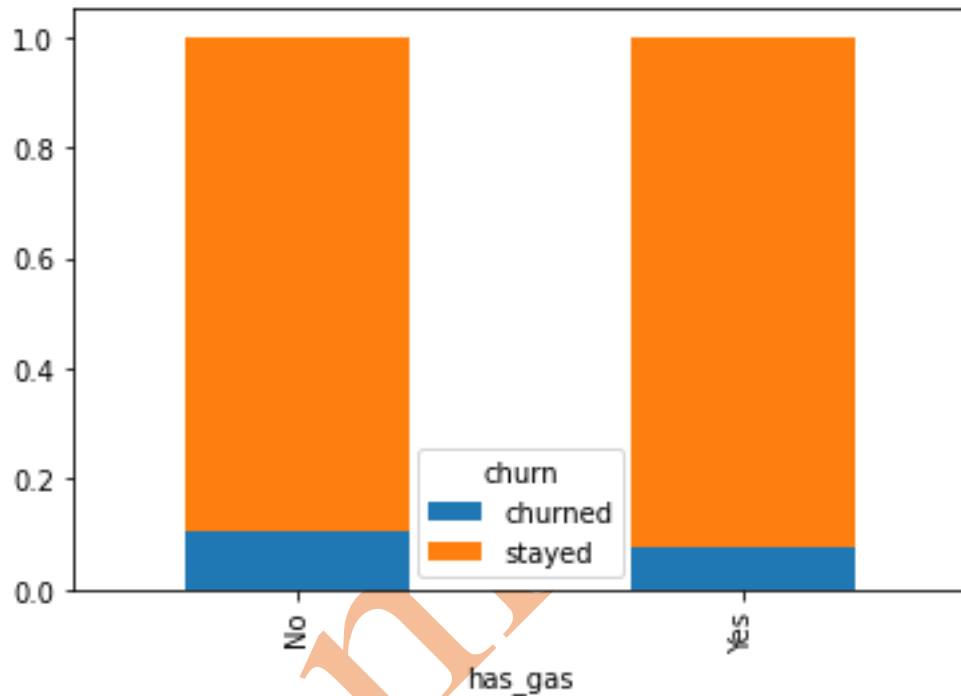


Relatively, companies with who joined the firm through ‘foosdfpfkusacimwkcsosbicdxkicaua’ and ‘usilxuppasemubllopkaafesmlibmsdf’ sales channel are more likely to churn.

Finding out if churning is dependent on whether a company has gas or not

```
[396]: pd.crosstab(merge['has_gas'],merge['churn'],normalize='index').plot.
        bar(stacked=True)
```

```
[396]: <AxesSubplot:xlabel='has_gas'>
```



It can be seen that, relatively, companies without gas churned more than the companies with gas. Finding out if churning is dependent on whether a company subscribed power or not

```
[397]: merge.groupby(['pow_max','churn']).size().unstack().head()
```

```
[397]: churn  churned  stayed
pow_max
1.000      NaN      1.0  3.300
2.0        5.0  3.450   2.0
2.0  3.464      NaN   4.0
3.500        1.0      1.0
```

```
[398]: pow_id = merge[['id','churn','pow_max']]
        pow_id.head()
```

```
[398]:          id  churn pow_max
0  48ada52261e7cf58715202705a0451c9  stayed  180.000
1  24011ae4ebbe3035111d65fa7c15bc57  churned   43.648
```

2	d29c2c54acc38ff3c0614d0a653813dd	stayed	13.800
3	764c75f661154dac3a6c254cd082ea7d	stayed	13.856
4	bba03439a292a1e166f80264c16191cb	stayed	13.200 we need to group

the pow\_max into grades

```
[399]: bins = [0,10,20,30,40,50,60,70,80,90,100,110,120,130,140,200]
```

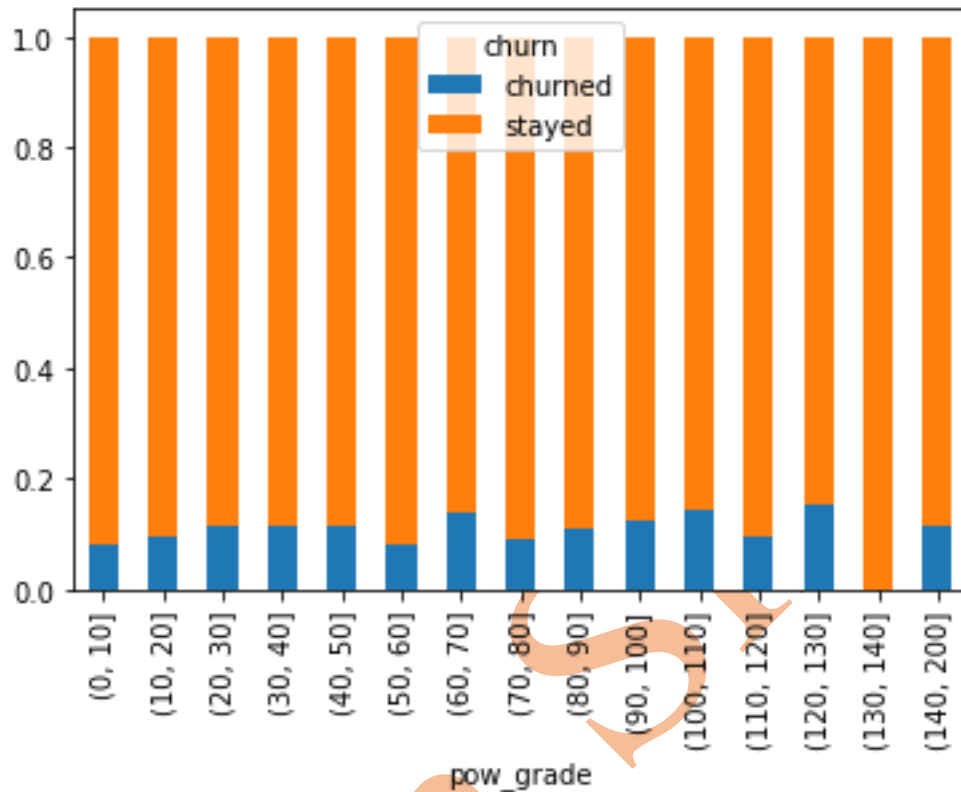
```
[400]: pow_id['pow_grade'] = pd.cut(pow_id['pow_max'],bins, labels=None)
pow_id.head()
```

```
[400]:
```

	id	churn	pow_max	pow_grade
0	48ada52261e7cf58715202705a0451c9	stayed	180.000	(140, 200]
1	24011ae4ebbe3035111d65fa7c15bc57	churned	43.648	(40, 50]
2	d29c2c54acc38ff3c0614d0a653813dd	stayed	13.800	(10, 20]
3	764c75f661154dac3a6c254cd082ea7d	stayed	13.856	(10, 20]
4	bba03439a292a1e166f80264c16191cb	stayed	13.200	(10, 20]

```
[401]: plt.figure(figsize = (40, 20))
pd.crosstab(pow_id['pow_grade'],pow_id['churn'],normalize='index').
plot. .bar(stacked=True)
```

```
[401]: <AxesSubplot:xlabel='pow_grade'>
<Figure size 2880x1440 with 0 Axes>
```



Relatively, companies with (60,70), (100,110) and (120,130) suscribed power are more likely to churn.

```
[402]: merge['date_activ'] =
pd.to_datetime(merge['date_activ']).dt.to_period('m')
merge['date_end'] =
pd.to_datetime(merge['date_end']).dt.to_period('m')
merge['date_modif_prod'] =
pd.to_datetime(merge['date_modif_prod']).dt.
.to_period('m') merge['date_renewal'] =
pd.to_datetime(merge['date_renewal']).dt.to_period('m')
```

```
[403]: merge.head()
```

```
[403]:      id      churn \
0  48ada52261e7cf58715202705a0451c9  stayed
1  24011ae4ebbe3035111d65fa7c15bc57  churned
2  d29c2c54acc38ff3c0614d0a653813dd  stayed
3  764c75f661154dac3a6c254cd082ea7d  stayed
4  bba03439a292a1e166f80264c16191cb  stayed
```

activity\_new

channel\_sales \



```

0          esoiifxdlbkcsluxmfuacbdckommixw
          lmkebamcaaclubfxadlmueccxoimlema
1          NaN foosdfpfkusacimwkcsosbicdxkicaua
2          NaN NaN
3          NaN foosdfpfkusacimwkcsosbicdxkicaua
4          NaN lmkebamcaaclubfxadlmueccxoimlema

cons_12m cons_gas_12m cons_last_month date_activ date_end \
0          309275 0          10025 2012-11 2016-11
1          0 54946 0          2013-06 2016-06
2          4660 0          0          2009-08 2016-08
3          544 0 0 2010-04 2016-04 4 1584 0 0 2010-03 2016-03

date_modif_prod ... forecast_price_pow_pl has_gas imp_cons \
0          2012-11 ... 58.995952 No 831.8 1 NaT ...
          40.606701 Yes 0.0
2          2009-08 ... 44.311378 No 0.0
3          2010-04 ... 44.311378 No 0.0
4          2010-03 ... 44.311378 No 0.0

margin_gross_pow_ele margin_net_pow_ele nb_prod_act net_margin \
0          -41.76 -41.76 1 1732.36
1          25.44 25.44 2 678.99
2          16.38 16.38 1 18.89
3          28.60 28.60 1 6.60 4 30.22 30.22 1 25.46

num_years_antig          origin_up pow_max
0          3 ldkssxwpmemidmecebumciepifcamkci 180.000
1          3 lxidpiddsbxsbosboudacockeimpuepw 43.648
2          6 kamkkxfxxuwbdsldkwifmmcsiusuosws 13.800
3          6 kamkkxfxxuwbdsldkwifmmcsiusuosws 13.856
4          6 kamkkxfxxuwbdsldkwifmmcsiusuosws 13.200

```

[5 rows x 27 columns]

```


[404]: plt.rcParams['figure.figsize']=(20,10)
merge.groupby(['date_activ','churn']).size().unstack().plot.bar(sta
cked=True)

```

```

[404]: <AxesSubplot:xlabel='date_activ'>


```



Relatively, companies with who joined the firm through from december 2011 to 2014 are more likely to churn.

```
[405]: plt.rcParams['figure.figsize']=(20,10)
merge.groupby(['date_end', 'churn']).size().unstack().plot.bar(stack
ed=True)
```


```
[405]: <AxesSubplot:xlabel='date_end'>
```



Active products and services

```
[406]: plt.rcParams['figure.figsize']=(10,5)
pd.crosstab(merge['nb_prod_act'],merge['churn'],normalize='index').
plot. .bar(stacked=True)
```


```
[406]: <AxesSubplot:xlabel='nb_prod_act'>
```



Relatively, companies with 1-5 active products and services are more likely to churn.

```
[407]: plt.rcParams['figure.figsize']=(10,5)
pd.crosstab(merge['num_years_antig'],merge['churn'],normalize='index').plot. .bar(stacked=True)
```

```
[407]: <AxesSubplot:xlabel='num_years_antig'>
```




Relatively, companies are likely to churn in the second year. After, the probability for companies to churn diminishes up to the ninth year and starts rising again. By the 15th year, the companies are more likely to stay.

Code of the electricity campaign the customer first subscribed to

```
[408]: plt.rcParams['figure.figsize']=(10,5)
pd.crosstab(merge['origin_up'],merge['churn'],normalize='index').plot.
bar(stacked=True)
```

```
[408]: <AxesSubplot:xlabel='origin_up'>
```




Relatively, companies that first subscribed to the code of the electricity campaign 'lxidpiddsbx' is likely to churn.

## 11 Histogram of the data


```
[409]: merge["cons_12m"].plot.hist()
```

```
[409]: <AxesSubplot:ylabel='Frequency'>
```




```
[410]: merge["cons_gas_12m"].plot.hist()
```

```
[410]: <AxesSubplot:ylabel='Frequency'>
```




```
[411]: merge["forecast_cons_12m"].plot.hist()
```

```
[411]: <AxesSubplot:ylabel='Frequency'>
```




```
[412]: merge["forecast_cons_year"].plot.hist()
```

```
[412]: <AxesSubplot:ylabel='Frequency'>
```




```
[413]: merge["forecast_discount_energy"].plot.hist()
```

```
[413]: <AxesSubplot:ylabel='Frequency'>
```




```
[414]: merge["forecast_meter_rent_12m"].plot.hist()
```

```
[414]: <AxesSubplot:ylabel='Frequency'>
```



```
[415]: merge["forecast_price_energy_p1"].plot.hist()
<AxesSubplot:ylabel='Frequency'>
```



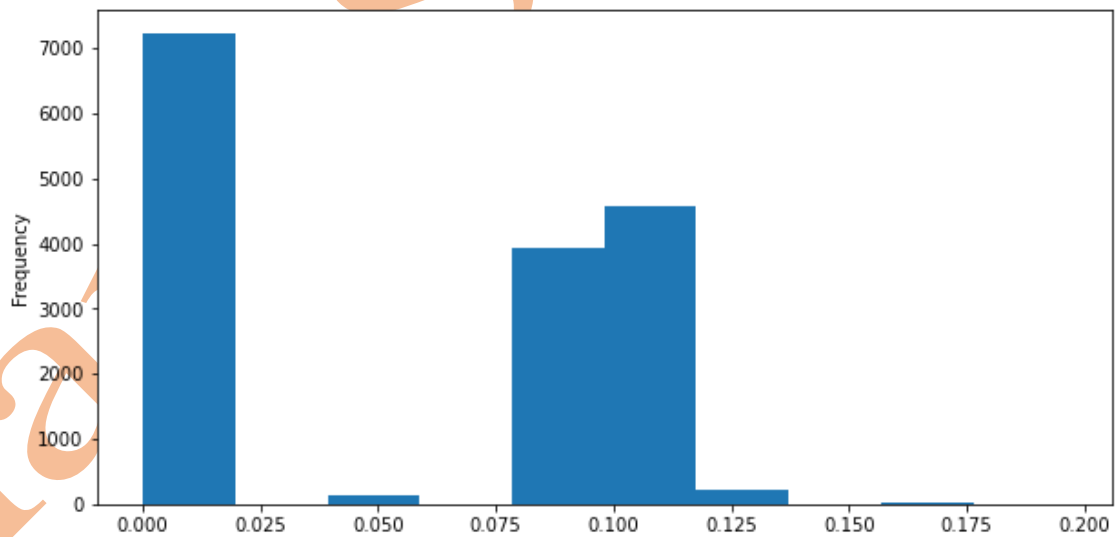



[415]:

[416]:

```
merge["forecast_price_energy_p2"].plot.hist()
```

[416]: <AxesSubplot:ylabel='Frequency'>






```
[417]: merge["forecast_price_pow_pl"].plot.hist()
```

```
[417]: <AxesSubplot:ylabel='Frequency'>
```

```
[418]: merge["imp_cons"].plot.hist()
```


```
[418]: <AxesSubplot:ylabel='Frequency'>
```



```
<AxesSubplot:ylabel='Frequency'>
```


```
[419]: merge["margin_gross_pow_ele"].plot.hist()
```

```
[419]:
```




```
[420]: merge["nb_prod_act"].plot.hist()
```

```
[420]: <AxesSubplot:ylabel='Frequency'>
```



```
[421]: merge["net_margin"].plot.hist()
```


```
<AxesSubplot:ylabel='Frequency'>
```




[421]:

```
[422]: merge["num_years_antig"].plot.hist()
```

[422]: <AxesSubplot:ylabel='Frequency'>





```
[423]: merge["pow_max"].plot.hist()
```

[423]:

It can be seen that most of the data are rightly skewed

### 11.1 Checking the company with the highest consumption

```
[424]: consumption = merge[["id", "cons_12m",
    "cons_gas_12m", "cons_last_month", "imp_cons", "has_gas", "churn"]]
```

```
[425]: total_cons_12m = pd.DataFrame(consumption.groupby(["id",
    "churn"])["cons_12m"].
    .agg(["sum"]))
```

```
total_cons_12m.sort_values(ascending=False,
    by="sum").head()
```

```
[425]: sum id      churn
2c2abbe8998364dd500e41588d41f45f      stayed
16097108      b880901f75613c801886354abf24f30a
stayed      6286272
3cbf266f90f0419636aa9e748fa0e7f0      stayed
<AxesSubplot:ylabel='Frequency'>
```

```
6286272    f3baf732b3a86a45f5aec2d4578070c0
stayed                                           6286272
4130bb214991c2ec4504b96d527624ca    stayed
6286272
```

It can be seen that ,company '2c2abbe8998364dd500e41588d41f45f' has the highest consumption of energy.

```
[426]: merge.to_csv('merge.csv')
```